

---

## ▼ Everything is Better with Friends

### Using SAS in Python Applications with SASPy and Open-Source Tooling (Beyond the Basics)

## ▼ Setup for Part 1

### Getting setup to use Google Colab with SAS OnDemand for Academics (ODA)

1. To execute code cells, you'll need credentials for the following accounts:

- Google. (If you're not already signed in, you should see a **Sign In** button in the upper right corner. You can also visit <https://accounts.google.com/signup> to create an account for free.)
- SAS OnDemand for Academics. (You can create an account for free at <https://welcome.oda.sas.com/> using an existing SAS Profile account. If you don't already have a SAS Profile account, you can create one for free using the "Don't have a SAS Profile?" link on the ODA login page.)

2. We recommend enabling line numbers using the Tools menu: **Tools** -> **Settings** -> **Editor** -> **Show line numbers** -> **Save**

3. We also recommend enabling the Table of Contents using the View menu: **View** -> **Table of contents**

4. To save a copy of this notebook, along with any edits you make, please use the File menu: **File** -> **Save a copy in Drive**

5. Looking for "extra credit"? Please let us know if you spot any typos!

## ▼ Connect to SAS OnDemand for Academics (ODA) and start a SAS session

### Instructions:

1. Determine the Region for your ODA account by logging into <https://welcome.oda.sas.com/>. You should see the value `Asia Pacific`, `Europe`, or `United States` next to your username in the upper-right corner. (For more information about Regions, please see the [ODA documentation](#).)
2. If your ODA account is associated with a Region other than `United States`, comment out Line 11 by adding a number sign (`#`) at the beginning of the line, and then do the following:
  - If your ODA account is associated with the Region `Europe`, uncomment Line 14 by removing the number sign (`#`) at the beginning of the line.
  - If your ODA account is associated with the Region `Asia Pacific`, uncomment Line 17 by removing the number sign (`#`) at the beginning of the line.
3. Click anywhere in the code cell, and run the cell using Shift-Enter.
4. At the prompt `Please enter the IOM user id`, enter either your SAS ODA user ID or the email address associated with your ODA account.
5. At the prompt `Please enter the password for IOM user`, enter the password for your SAS ODA account.

```
1 !pip install saspy
2
3 import saspy
4
5 sas = saspy.SASsession(
6     java='/usr/bin/java',
7     iomport=8591,
8     encoding='utf-8',
9
10    # The following line should be uncommented if, and only if, your ODA account is associated with the Region
11    iomhost = ['odaws01-usw2.oda.sas.com', 'odaws02-usw2.oda.sas.com', 'odaws03-usw2.oda.sas.com', 'odaws04-usw
12
13    # The following line should be uncommented if, and only if, your ODA account is associated with the Region
14    #iomhost = ['odaws01-euw1.oda.sas.com', 'odaws02-euw1.oda.sas.com'],
15
16    # The following line should be uncommented if, and only if, your ODA account is associated with the Region
17    #iomhost = ['odaws01-apse1.oda.sas.com', 'odaws02-apse1.oda.sas.com'],
18
```

```
19 )
20 print(sas)
```

```
Collecting saspy
  Downloading saspy-4.3.0.tar.gz (9.9 MB)
    |██████████████████████████████████████| 9.9 MB 5.5 MB/s
Building wheels for collected packages: saspy
  Building wheel for saspy (setup.py) ... done
  Created wheel for saspy: filename=saspy-4.3.0-py3-none-any.whl size=9929656 sha256=1c8815f0177820993517a4d:
  Stored in directory: /root/.cache/pip/wheels/c3/b5/08/62c85da319a5178d19559f996ceefd7583b9bf31feeafbad8e
Successfully built saspy
Installing collected packages: saspy
Successfully installed saspy-4.3.0
Using SAS Config named: default
Please enter the IOM user id: isaiah.lankham@ucop.edu
Please enter the password for IOM user : .....
SAS Connection established. Subprocess id is 135

Access Method          = IOM
SAS Config name        = default
SAS Config file        = /usr/local/lib/python3.7/dist-packages/saspy/sascfg.py
WORK Path              = /saswork/SAS_work3160000013E6_odaws03-usw2.oda.sas.com/SAS_workB05A000013E6_odaws03-1
SAS Version            = 9.04.01M6P11072018
SASPy Version          = 4.3.0
Teach me SAS           = False
Batch                  = False
Results                = Pandas
SAS Session Encoding   = utf-8
Python Encoding value  = utf-8
SAS process Pid value  = 5094
```

**Note:** This establishes a connection from Python in Google Colab to a SAS session running in SAS ODA.

## ▼ Install and import additional packages

```

1 # Install the rich module for colorful printing
2 !pip install rich
3
4 # We'll use IPython to display DataFrames or HTML content
5 from IPython.display import display, HTML
6
7 # We're overwriting the default print function with rich.print
8 from rich import print
9
10 # We're also setting the maximum line width of rich.print to be a bit wider (to avoid line wrapping)
11 from rich import get_console
12 console = get_console()
13 console.width = 165

```

Collecting rich

Downloading rich-12.4.1-py3-none-any.whl (231 kB)

|██| 231 kB 4.7 MB/s

Requirement already satisfied: pygments<3.0.0,>=2.6.0 in /usr/local/lib/python3.7/dist-packages (from rich)

Collecting commonmark<0.10.0,>=0.9.0

Downloading commonmark-0.9.1-py2.py3-none-any.whl (51 kB)

|██| 51 kB 4.8 MB/s

Requirement already satisfied: typing-extensions<5.0,>=4.0.0 in /usr/local/lib/python3.7/dist-packages (from

Installing collected packages: commonmark, rich

Successfully installed commonmark-0.9.1 rich-12.4.1

## ▼ Part 1. Calling SAS/STAT procedures in Python applications

### ▼ Section 1.1. Create titanic\_sds

```

1 # I'd rather not have spaces and slashes in my SAS variable names.
2 sas.submit('options validvarname=v7;')
3
4 # Read the titanic dataset into SAS from a URL.
5 titanic_url = 'https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/stuff/titanic.csv'

```

```
6 titanic_sds = sas.read_csv(file=titanic_url,table='titanic',libref='work')
7
8 # What kind of object is titanic_sds?
9 print(type(titanic_sds))
10 print('\n')
11 print(titanic_sds)
12
13 # Curious what's under the hood?
14 print('Here\'s the SAS code being used by SASPy to import the dataset:')
15 sas.teach_me_SAS(True)
16 sas.read_csv(file=titanic_url,table='titanic',libref='work')
17 sas.teach_me_SAS(False)
18 print('\n')
19
20 # Let's take a look at the data.
21 display(titanic_sds.columnInfo())
22 display(titanic_sds.head())
```

```
<class 'saspy.sasdata.SASdata'>
```

```
Libref = work  
Table = titanic  
Dsopts = {}  
Results = Pandas
```

Here's the SAS code being used by SASPy to import the dataset:

```
filename x url "https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/  
proc import datafile=x out=work.'titanic' dbms=csv replace; run;
```

	Member	Num	Variable	Type	Len	Pos	Format	Informa
0	WORK.TITANIC	5.0	Age	Num	8.0	16.0	BEST12.	BEST3
1	WORK.TITANIC	8.0	Fare	Num	8.0	40.0	BEST12.	BEST3

### Concept Check 1.1

- Try this, and see what happens: Comment out Line 17 so that `sas.teach_me_SAS` is no longer turned off, and then rerun the code cell above.
- True or False: Commenting out Line 17, so that `sas.teach_me_SAS` is left on, has no effect on subsequent code.
- Fun Fact: The SASdata object `titanic_sds` does not represent a dataset kept in memory in our local Google Colab session. Instead, it's a pointer to normal SAS dataset file kept on disk in the remote SAS ODA session.

**Solution:** False! If we don't turn "Teach Me SAS" off, Lines 21-22 won't execute.

```
0      0.0      3.0      Harris      male      22.0      1.0
```

### ▼ Section 1.2. Create titanic\_partitions

Bradley

```
1 # Make sure sas.teach_me_SAS is set to False.  
2 sas.teach_me_SAS(False)  
3  
4 # Let's partition the dataset into subsets for training and testing a predictive model.  
5 titanic_partitions = titanic_sds.partition(seed=42, singleOut=False)  
6  
7 # And then print some information about each subset.
```

```

8  print(type(titanic_partitions))
9  print('\n')
10 print(titanic_partitions)
11 print('\n')
12 print(f'{titanic_partitions[0].obs()} observations for training')
13 print(f'{titanic_partitions[1].obs()} observations for test')

```

```
<class 'tuple'>
```

```

(Libref = work
Table   = titanic1_train
Dsopts  =
Results = Pandas
, Libref = work
Table   = titanic1_score
Dsopts  =
Results = Pandas
)

```

```
621 observations for training
```

```
266 observations for test
```

## Concept Check 1.2

- Try this, and see what happens: Try using `sas.teach_me_SAS` to see the SAS code generated by the `partition` method on Line 5.
- True or False: In order to make sure `titanic_partitions` is created as expected (without a Traceback, a.k.a., a run-time error), Line 5 needs to be run with `sas.teach_me_SAS` turned off.
- Fun Fact: A `tuple` is like a fixed-length `list`. Whereas a `list` can be changed in place, and can grow and shrink in size, a `tuple` is immutable. This means any attempt to change a `tuple` actually results in a new `tuple` being created, instead. (Other examples of immutable Python objects include `int` and `str`.)

**Solution:** False! If we don't turn "Teach Me SAS" off, Lines 8-13 will result in Traceback errors since the variable `titanic_partitions` won't have been created.

### ▼ Section 1.3. Create titanic\_model

```
1 # Make sure sas.teach_me_SAS is set to False (just in case).
2 sas.teach_me_SAS(False)
3
4 # Now let's create an object that will allow us to access SAS/STAT procedures.
5 sas_stat = sas.sasstat()
6
7 # We'll also set our response variables and explanatory variables.
8 outcome = "survived(event='1')"
9 covariates = 'pclass sex age siblings_spouses_aboard parents_children_aboard fare'
10
11 # We're now ready to use PROC LOGISTIC to train a model and score the test dataset.
12 titanic_model = sas_stat.logistic(
13     data = titanic_partitions[0],
14     cls = 'sex',
15     model = f'{outcome} = {covariates}',
16     stmtpassthrough = f'score data={titanic_partitions[1].table} out=scored fitstat',
17     procopts = 'plots=none',
18 )
19
20 # Let's check the SAS log to see how everything went.
21 display(titanic_model.LOG)
22
23 # What else is in this titanic_model object?
24 print(titanic_model.__dict__['_names'])
```



63

The SAS System

408        options nosource;

536

64

The SAS System

537

538        %macro proccall(d);

539        proc logistic data=work.'titanic1\_train'n   plot=all plots=none ;

540        class sex;

541        model survived(event='1') = pclass sex age siblings\_spouses\_aboa

542        score data=titanic1\_score out=scored fitstat ;

543        run; quit; %mend;

544        %mangobj(log0001,logistic,'titanic1\_train'n);

548

549

550

65

The SAS System

551

[

```
'ASSOCIATION',
'CLASSLEVELTME0'
```

### Concept Check 1.3

- Try this, and see what happens: Change `LOG` on Line 21 to any of the object names in the list that was output by Line 24.
- True or False: The Python object `titanic_model` has 13 sub-objects nested inside of it.
- Fun Fact: The strings `"survived(event='1')"` and `'survived(event="1")'` produce identical behavior in Python.

```
'RESPONSENOTICE',
'LOG0001FITSTAT'
```

**Solution:** True! There are 13 items in the list printed out by Line 24.

### ▼ Section 1.4. Create `train_auc` and `test_auc`

```

1 # But how does the model perform?
2 display(titanic_model.ASSOCIATION)
3 display(titanic_model.SCOREFITSTAT)
4
5 # Let's pull out the AUC value for training.
6 train_auc_srs = titanic_model.ASSOCIATION.loc[titanic_model.ASSOCIATION['Label2'] == 'c', 'nValue2']
7 print('\n')
8 print(type(train_auc_srs))
9 print('\n')
10 print(train_auc_srs)
11 train_auc = train_auc_srs.iloc[0]
12
13 # Let's also pull out the AUC value for test.
14 test_auc = titanic_model.SCOREFITSTAT['AUC'][0]
15
16 # And, finally, let's compare them.
17 print('\n')
18 print(f'Training AUC: {train_auc:.4f}, Test AUC: {test_auc:.4f}')

```

	Label1	cValue1	nValue1	Label2	cValue2	nValue2	
0	Percent Concordant	86.5	86.532990	Somers' D	0.731	0.730761	
1	Percent Discordant	13.5	13.456928	Gamma	0.731	0.730834	
2	Percent Tied	0.0	0.010082	Tau-a	0.339	0.338866	
3	Pairs	89270	89270.000000	c	0.865	0.865380	
	DataSet	Freq	LogLike	MisClass	AUC	AIC	A
0	WORK.TITANIC1_SCORE	266.0	-128.118646	0.184211	0.841264	270.237292	270.6

```
<class 'pandas.core.series.Series'>
```

```
3      0.86538
```

```
Name: nValue2, dtype: float64
```

```
Training AUC: 0.8654, Test AUC: 0.8413
```

## Concept Check 1.4

- Short Answer: How would you access specific cells of a dataset in SAS?
- Fun Fact: A DataFrame in Python is a tabular data structure with rows and columns, similar to a SAS data set. However, while SAS datasets are typically accessed from disk and processed row-by-row, DataFrames are loaded into memory all at once. This means values in DataFrames can be randomly accessed (like on Line 6 above), but it also means the size of DataFrames can't grow beyond available memory. Problem solving always involves trade-offs. Python provides flexible access to in-memory data, whereas SAS makes it possible to work with datasets whose size exceeds available memory.

**Solution:** We can subset a SAS dataset to specific rows and columns with KEEP, DROP, and WHERE dataset options, or in a SQL query with SELECT, WHERE, and HAVING clauses. To pull out a single value, we might store it in a macro variable with CALL SYMPUTX in a DATA step, or an INTO expression in PROC SQL.

## ▼ Section 1.5. Alternate ways to create train\_auc

```
1 # We could simplify getting the training AUC by just hard-coding its row index.
2 train_auc = titanic_model.ASSOCIATION['nValue2'][3]
3 print(train_auc)
```

**0.86538030693402**

```
1 # Or we can create an index on the Label2 column and get the best of both worlds.
2 indexed_association = titanic_model.ASSOCIATION.set_index('Label2')
3 train_auc = indexed_association['nValue2']['c']
4 print(train_auc)
```

**0.86538030693402**

## Concept Check 1.5

- Multiple choice: What do you think is the best way to pull out the AUC value from the ASSOCIATION attribute of `titanic_model`?
  - A. Using `loc` to subset on the value of a column.
  - B. Using column labels and integer row indices.
  - C. Using `set_index` to create a non-integer index.
- Fun Fact: This exercise would seem to contradict the often-quoted aphorism "There should be one-- and preferably only one -- obvious way to do it" from the [Zen of Python](#)

**Solution:** The authors prefers option C (`set_index`), but there are merits to all three methods.

## ▼ Section 1.6. Additional Exercises

For practice, we recommend the following:

1. Run the code in the cell below (taken from Section 1.4).
2. Modify the code to instead pull the `Percent Concordant` value out of `titanic_model.ASSOCIATION`.
3. Use the `display` method to view the contents of `titanic_model.PARAMETERESTIMATES`.
4. Finally, pull a single value out of `titanic_model.PARAMETERESTIMATES`. (E.g., you might try getting the parameter estimate for the variable `Age`, which you can do using any of the three methods discussed above.)

```
1 display(titanic_model.ASSOCIATION)
2 train_percent_concordant_srs = titanic_model.ASSOCIATION.loc[titanic_model.ASSOCIATION['Label1'] == 'Percent Co
3 print(train_percent_concordant_srs)
4
5 display(titanic_model.PARAMETERESTIMATES)
6 age_parameter_estimate = titanic_model.PARAMETERESTIMATES.loc[titanic_model.PARAMETERESTIMATES['Variable'] == '
7 print(age_parameter_estimate)
```

	Label1	cValue1	nValue1	Label2	cValue2	nValue2
0	Percent Concordant	86.5	86.532990	Somers' D	0.731	0.730761
1	Percent Discordant	13.5	13.456928	Gamma	0.731	0.730834
2	Percent Tied	0.0	0.010082	Tau-a	0.339	0.338866
3	Pairs	89270	89270.000000	c	0.865	0.865380

0 86.53299

Name: nValue1, dtype: float64

	Variable	ClassVal0	DF	Estimate	StdErr	WaldChiSq	ProbChiSq
0	Intercept	NaN	1.0	4.284323	0.647637	43.762334	3.707725e-11
1	Pclass	NaN	1.0	-1.345140	0.180891	55.297173	1.036175e-13
2	Sex	female	1.0	1.441836	0.125068	132.903631	9.490878e-31
3	Age	NaN	1.0	-0.044455	0.009402	22.354893	2.266349e-06
4	Siblings Spouses Abn	NaN	1.0	-0.372227	0.137423	7.336604	6.756436e-02

## ▼ Notes and Resources

1. For more about the `pandas` package, including the methods used above, see the following:

- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.loc.html>
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.iloc.html>
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html>
- [https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.set\\_index.html](https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.set_index.html)
- <https://pandas.pydata.org/docs/reference/api/pandas.Series.html>

2. For more about the `rich` package, see <https://rich.readthedocs.io/>
3. For more about the `saspy` package, including the methods used above, see the following:
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasdata.SASdata.columnInfo>
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasdata.SASdata.head>
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasdata.SASdata.obs>
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasdata.SASdata.partition>
  - <https://sassoftware.github.io/saspy/api.html#saspy.SASsession.lastlog>
  - [https://sassoftware.github.io/saspy/api.html#saspy.SASsession.read\\_csv](https://sassoftware.github.io/saspy/api.html#saspy.SASsession.read_csv)
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasstat.SASstat>
  - <https://sassoftware.github.io/saspy/api.html#saspy.sasstat.SASstat.logistic>
4. For more information on built-in Python data structures, such as tuples, see <https://jakevdp.github.io/WhirlwindTourOfPython/06-built-in-data-structures.html>.
5. For more information on f-strings (i.e., Python strings like `f'{outcome} = {covariates}'`), see <https://realpython.com/python-f-strings/>.
6. We welcome follow-up conversations. You can connect with us on LinkedIn or email us at [isaiah.lankham@gmail.com](mailto:isaiah.lankham@gmail.com) and [matthew.t.slaughter@gmail.com](mailto:matthew.t.slaughter@gmail.com)
7. If you have a GitHub account (or don't mind creating one), you can also chat with us on Gitter at <https://gitter.im/saspy-bffs/community>